

Know Your Body: A Harness for Direct and Self-Improving Robot Control with VLMs

<https://loule0-0.github.io/KnowBody/>

Zeyu Lou^{1,2} Yanhong Zeng² Yong Wang^{2,3} Chenyang Si^{1,†}

¹Nanjing University ²Ant Group ³Zhejiang University

Abstract

A general-purpose vision-language model can understand a task goal without knowing how a particular robot’s motion and functional parts produce the intended effect. We introduce **KnowBody**, a harness that makes these action-relevant body relations explicit, queryable, and revisable while keeping the model weights frozen. Initialized from one off-task trajectory, a partial body model guides action selection and the interpretation of past interactions. New evidence refines the model, and knowledge dependent on revised body estimates is rechecked before reuse. Across 32 fixed-budget trials on four real-robot tasks, initialized KnowBody achieves 75% completion versus 25% for the native harness and requires fewer planner rounds on successful trials in tasks completed by both. With persistent updates enabled, planner rounds decrease by 29–53% from the first to the fifth recorded success.

1 Introduction

General-purpose vision-language models (VLMs) offer an appealing route to robot manipulation: their broad visual and task knowledge can support decisions across tasks, while a harness connects those decisions to robot observations, action tools, and execution feedback. This approach makes it possible to attempt new tasks without first training a task-specific policy, and offers a way to bring advances in general-purpose models into robot control. Yet broad task knowledge does not establish how a particular robot’s motions and functional parts produce the intended physical effects.

Consider writing with a grasped pen: the task requires contact at the pen tip, whereas the robot receives end-effector pose commands. Choosing those commands requires knowing how the held tip relates to the commanded pose. Observations and interaction histories provide evidence about such relations; using that evidence for control requires identifying what is supported, applying it to the next action, and revising it when new observations reveal a mismatch. The central challenge is therefore to connect action selection and learn-




Figure 1. Knowing the task is not knowing the body. A VLM is able to recognize what to write, but the robot receives an end-effector pose while writing acts at the held tip. KNOWBODY makes this command-effect relation explicit for action decision and outcomes interpretation.

ing through explicit body knowledge that a general-purpose VLM can use and that subsequent interactions can test and refine.

We introduce **KnowBody**, a harness that connects a VLM’s action decisions to a partial, queryable body model. Before acting, the VLM uses relevant motion and functional-geometry relations, together with current observations and knowledge of previous action outcomes, to choose concrete robot commands. Execution feedback and subsequent observations then serve two purposes: assessing task progress and testing the body relations used in the decision. Evidence from these interactions supports the refinement of body estimates and the accumulation of conditional knowledge about action outcomes, which inform later decisions. The body model thus provides a shared reference for choosing actions and learning from their consequences, while the VLM remains responsible for task reasoning and action selection.

A single off-task trajectory initializes the body model with relations supported by its recorded actions and observations. Subsequent interactions provide evidence for refining body

[†]Corresponding author: Chenyang Si. Emails: zeyu.lou@mail@gmail.com, zengyh1900@gmail.com, yongw@zju.edu.cn, and chenyang.si@nju.edu.cn.




Figure 2. KnowBody separates online state feedback from validated cross-episode revision. At each decision, the VLM receives the task, current observations, current state M , eligible rules K , and selected original evidence E . It queries the partial body model, issues waypoints and gripper commands, and receives execution feedback through Observe → Record → Check → Update M ; only M returns directly to the next decision. Qualified evidence enters the revision lane, where candidate B/K changes are independently validated and rules dependent on a changed B are re-evaluated. Dashed purple arrows publish validated B_{k+1} to the body model and eligible K_{k+1} to the decision context for later episodes. One off-task trajectory initializes supported relations in B_0 .

estimates and acquiring conditional knowledge about action outcomes. Validated updates are retained across episodes, allowing later decisions to draw on knowledge accumulated during earlier tasks. When a body estimate changes, interaction knowledge that depends on it is rechecked before reuse.

We evaluate KnowBody on four real-robot tasks: duck placement, apple pushing, writing, and pouring. Across 32 fixed-budget trials with the same frozen VLM and cross-episode updates disabled, KnowBody achieves 75% completion versus 25% for the native Codex harness. On the three tasks completed by both methods, successful KnowBody trials require 17–44% fewer planner rounds on average. With persistent updates enabled, rounds decrease by 29–53% from the first to the fifth recorded success across the four tasks.

2 Related Work

General models as robot decision makers. SayCan grounds language-model proposals with skill-value estimates [1], while PaLM-E incorporates continuous sensor observations into a trained embodied language model [13]. Code as Policies generates programs over perception and control APIs [27]; CaP-X examines how primitive abstraction, visual grounding, and execution feedback affect coding-agent manipulation [14]. PIVOT iteratively selects image-space actions [30], and VLMPC uses VLM-proposed sequences in a predictive-control loop [38].

Show-Harness interprets discrete semantic motion units as local robot actions, with multi-view guidance, proprioception, and action chunking [7]. Agent as Policy (AGP) uses calibrated observations, geometric queries, and runtime programs to generate and revise motion [21]. GPT-Policy compiles demonstrations, goal images, interaction history, and execution feedback into context for a fixed VLM acting

through a constrained controller [8]. KNOWBODY builds on these direct-control interfaces, focusing on body evidence from one off-task trajectory and on experience whose validity depends on that evidence.

Trained visuomotor policies and VLAs. Diffusion Policy learns multimodal action sequences from robot demonstrations [9], and RT-1 studies how model and data scale affect a multi-task real-robot policy [5]. RT-2 co-finetunes VLMs on robot trajectories and web vision-language data with actions represented as tokens [39]; OpenVLA similarly provides a pretrained, adaptable vision-language-action policy [22]. Reflective VLA trains an action expert to condition on observation–action–consequence histories, exposing deployment-specific actuation and calibration effects [26]. Harness VLA instead keeps a VLA frozen as a retryable contact primitive and uses execution memories to delimit the operating range of a fixed primitive library [36]. KNOWBODY instead controls the robot by interfacing a general VLM through a harness. Our matched evaluation therefore compares harnesses under the same VLM setting rather than trained VLA policies.

Explicit geometry and demonstration-conditioned transfer. Task and motion planning couples task-level choices with continuous geometric and motion constraints [15]. With foundation models, VoxPoser constructs 3D value maps [17], ReKep generates relational keypoint constraints [19], CoPa specifies grasp-part and post-grasp pose constraints [16], and RoboPoint predicts image-space affordance points [35]. These methods motivate explicit spatial structure for action generation.

Demonstrations provide another route: KAT and RoboPrompt encode examples for in-context action prediction [12, 34]; Vid2Robot learns from paired human videos and robot trajectories [20]; XSkill learns cross-embodiment skill prototypes [33]. UMI combines hand-held gripper demonstrations, relative action trajectories, and latency matching to train transferable robot policies [10]. Our off-task trajectory instead initializes supported motion correspondence, functional-part evidence, and camera context. It supplies body knowledge without prescribing the new task’s sequence or training a visuomotor policy.

Body models and observation. Continuous self-modeling enables recovery after morphological damage [3], and self-perception can recover a body scheme without a supplied kinematic model [32]. Later models learn forward behavior from action–sensation data [23] or pose-conditioned occupancy for planning and damage detection [6]. Human-CLAW studies VLM body reasoning in simulated egocentric tasks, separating action decisions from balance and motor-tracking failures [25]. Our body representation is partial: supported motion directions, functional points, temporary tool geometry, and observation context. Multi-view guidance in Show-Harness and iterative visual queries in PIVOT

inform observation selection [7, 30]; we also consider how a proposed action affects subsequent visibility.

Experience, memory, and validity under change. Inner Monologue feeds observations and success feedback into language-model planning [18]; REFLECT summarizes multisensory robot histories to explain failures and guide correction [28]. Reflexion retains verbal feedback [31], and ExpeL retrieves insights extracted from prior tasks [37]. RoboCat incorporates self-generated trajectories into subsequent training [4], whereas PhysMem verifies interaction hypotheses before reuse [24]. ASPIRE diagnoses failures from execution traces, validates repairs through re-execution, and consolidates reusable fixes into a growing skill library [29]. AGP retains procedures and corrections [21], while GPT-Policy adapts through demonstration and interaction context [8]. Relational transition models also support planning across object instances and configurations [11].

These works establish feedback, experience reuse, and validation as shared foundations. KNOWBODY addresses a specific dependency within this process: correcting body geometry can change the conditions inferred from an unchanged historical interaction. It therefore recomputes those conditions from the original observations and revalidates the affected rules before reuse. The distinction concerns how experience is maintained when its geometric basis changes, rather than the existence of memory or a verification loop.

3 Problem and Learning Protocol

At model–tool decision i , a general-purpose VLM receives instruction ℓ , current images and available depth o_t , measured pose $T_t = (R_t, p_t)$, gripper state, and available joint feedback. The index t advances with fresh observations and physical transitions, not with every read-only query. Declared robot and motion limits are supplied as interface constraints. The output u_i is a query, observation request, motion proposal, or termination request. A motion proposal specifies a chunk A_t of numerical Cartesian segments and optional gripper commands; a local executor performs kinematics and timing and records the achieved movement.

One continuous off-task episode initializes the body model:

$$\tau_0 = \{(o_m, a_m^{\text{req}}, a_m^{\text{ach}}, o_{m+1})\}_{m=1}^L, \quad B_0 = \mathcal{C}(\tau_0, \mathcal{H}), \quad (1)$$

Here $\mathcal{H}$ contains declared hardware metadata. The trajectory seeds supported motion–response relations by pairing observations with requested and achieved actions; independently versioned geometric measurements populate the corresponding functional-geometry components. A body snapshot B_k stores supported motion response, functional geometry, camera bindings, and component versions. Knowledge K_k stores conditional outcomes and their body dependencies; M_t retains temporary episode facts such as tool attachment. The




Figure 3. Action-relevant relations exposed by the body model. (a) Achieved robot motion is paired with the displacement of stationary landmarks in the wrist image, supporting visual-direction queries within the observed domain. (b) Permanent end-effector geometry is composed with the current attachment to map the commanded frame to the acting tool tip. The blue glyph schematically denotes the declared `panda_link8` frame; the green point marks the visible writing tip. (c) For a distant target, a valid endpoint solution does not certify the intervening motion. The translucent cutout anchors the current physical embodiment; the overlaid paths and body poses schematically contrast an endpoint-only transition with a discontinuous intermediate IK branch (red) against sequentially seeded waypoints that preserve a continuous full-arm path (green) to the same target. The query evaluates path feasibility using the kinematic relations in B and the current joint and attachment state in M before execution.

policy interface is

$$u_i \sim \pi_\theta(\ell, o_t, \mathcal{R}(B_k, K_k; o_t), M_t), \quad (2)$$

where $\mathcal{R}$ retrieves applicable evidence and model weights θ remain fixed. Adaptation changes B_k and K_k , not the VLM. Raw events remain immutable, and a fixed evaluation block cannot use its own outcomes to update its snapshot. Paused continuations of one physical episode remain one statistical group. The fixed-budget comparison and continual sequences in Section 5 evaluate task completion and repeated-task decision effort separately.

4 Method

KNOWBODY organizes robot interaction around a queryable, partial body model. Before acting, the model relates motion and functional-part geometry to the commands chosen by a frozen VLM. After acting, it supports checks of the observed response and supplies the geometric conditions under which experience is interpreted. These two roles couple body estimation to experience reuse: correcting a body estimate can change an earlier interaction’s conditions without changing its observed outcome. We describe the representation and initialization, body-conditioned interaction, evidence interpretation, and dependency-aware revision.

4.1 Body Representation and Initialization

A partial, queryable body model. We define body knowledge as action-relevant relations supported by observations and accessible through queries. Figure 3 previews three complementary forms in the order used for action reasoning: a local visual response relates achieved body motion to wrist-image displacement and supports inverse direction queries; permanent end-effector geometry combines with the current attachment to locate the active tip or part; and full-arm kinematic information distinguishes endpoint reachability from a continuous feasible path. In the last case, B provides the declared chain and its supported kinematic relations, while M provides the measured joint state and current attachment. Their query result conditions planner choices; the executor rechecks the selected motion against fresh state and retains final authority. A versioned body is

$$B_k = (C_k, G_k, R_k, D_k, V_k). \quad (3)$$

Here C_k identifies the physical configuration, declared kinematic chain, and camera acquisition modes; G_k represents permanent functional geometry; R_k represents the identifiable local motion–observation response; D_k specifies the supported operating domain; and V_k records sources and validation. A query returns a value with its frame, units, support, and available uncertainty. Relations outside the evidence-supported domain return unknown. The representation is therefore partial: it exposes identifiable relations

rather than assuming a complete dynamics model.

Separate body, state, and experience. Evidence $\mathcal{E}$ preserves observations, achieved motion, and attributed outcome records. Working state M_t describes the current scene, progress, and temporary attachments. Knowledge K_k stores validated conditional action–effect relations together with their supporting evidence and body dependencies. These stores have different update rules: new observations revise current state, whereas reusable body estimates and experience require validation.

For a held marker, permanent finger geometry belongs to G_k , while the grip-to-tip relation belongs to M_t^{attach} . Both contribute to a functional-part query, but regrasping invalidates the attachment rather than the permanent finger model. Each episode uses a fixed body version B_{pin} and a compatible knowledge snapshot while updating M_t . Contradictory evidence can suspend an affected query or rule during the episode; a replacement is activated only after validation.

Initialize from one off-task trajectory. The trajectory supplies observation–motion pairs, not a task-specific action sequence to imitate. For successive measured end-effector poses (R_m, p_m) and (R_{m+1}, p_{m+1}) , the achieved translation in the initial end-effector frame is $d_m^{\text{ef}} = R_m^\top (p_{m+1} - p_m)$. We pair this displacement with its actual before/after observations to initialize identifiable relations in B_0 . Each component records whether it derives from the seed trajectory, declared hardware, an independent measurement, or a later interaction.

The trajectory initializes a local visual response. For a rigid wrist camera, static scene points, and near-pure translation, normalized image correspondences $\xi, \xi' \in \mathbb{R}^2$ satisfy

$$Z'(\xi' - \xi) \approx J_B(\xi)d^{\text{ef}}, \quad (4)$$

where Z' is observed after-motion depth and $J_B(\xi) \in \mathbb{R}^{2 \times 3}$ is the fitted response. We reject inconsistent correspondences and fit within the translation subspace excited by the trajectory, balancing physical transitions so that dense image tracks do not dominate. This relation supports retrospective response checks and the direction query below; independently sourced measurements supply metric functional geometry. The parameterization and estimation procedure are given in [Appendix E.1](#).

4.2 Body-Conditioned Interaction

Assemble evidence for the next decision. At each decision, the harness supplies the VLM with the task goal and completion criteria, current images and measured robot state, working state M_t , relevant body-query results, and applicable knowledge. Selected records from $\mathcal{E}$ provide historical context when needed. Rule applicability is checked before relevance: a rule must match the current configuration, body dependencies, and observed conditions. Undetermined

conditions identify evidence to acquire, while unresolved counterevidence remains visible.

The VLM can query a body relation, request an observation, propose motion, or request termination. Tool results enter subsequent model inputs, allowing several queries before a physical action. After execution, the next decision is conditioned on achieved motion and its observations. The VLM weights remain frozen throughout.

Query a motion direction. The local visual response connects a desired image direction to robot translation. Let $s \in \mathbb{R}^2$ be the desired image direction at ξ , and let the columns of U form an orthonormal basis for the excited translation subspace. The query solves

$$d_0 = U[J_B(\xi)U]^\dagger s, \quad \hat{d}^{\text{base}} = R_t \frac{d_0}{\|d_0\|_2}, \quad (5)$$

where $\dagger$ denotes the pseudoinverse and R_t maps end-effector coordinates to the base frame. Positive depth scales image displacement without changing its direction, so this query does not require future depth. It returns a direction when support, conditioning, and solution checks pass. The VLM selects travel distance and orientation from the task and current observations; the query supplies a motion relation rather than a task policy.

Convert a functional target into a control target. Manipulation often acts through a fingertip, finger surface, or tool edge rather than the end-effector origin. For a functional part j with a supported point model,

$$x_j = p_t + R_t r_j(g_t), \quad (6)$$

where (R_t, p_t) is the measured end-effector pose, g_t is gripper state, and r_j is the end-effector-frame offset determined by B_{pin} and the current attachment. Given a target x_o and unit approach direction n in the same base frame, the query computes

$$\begin{aligned} \delta &= x_o - x_j, & h &= n^\top \delta, \\ e_\perp &= \|(I - nn^\top)\delta\|_2. \end{aligned} \quad (7)$$

The axial gap h and lateral offset $e_\perp$ distinguish approaching the target from aligning with it. Edge and surface interactions use their corresponding supported geometry rather than treating a point as the full interaction region.

For a desired part position x_j^* , orientation R^* , and gripper state g^* selected by the VLM, the control target is

$$p^* = x_j^* - R^* r_j(g^*). \quad (8)$$

This compensates for the offset between the commanded origin and the part producing the effect. Geometric queries require measured offsets and coordinate transforms. When uncertainty prevents a relation from determining a condition, the condition remains unknown and can motivate further observation.

Expose full-arm feasibility. A target pose is not a complete motion plan. For a sufficiently long or multi-waypoint proposal, a body query first solves the target pose and then evaluates the intervening route. Within each proposed segment it samples position linearly and orientation by quaternion interpolation, solves each pose from the preceding joint solution, and checks the resulting full-arm path. The query reports endpoint reachability separately from continuous-path feasibility, including the first failed stage when the latter is rejected. It is bound to the pinned kinematic description in B and the measured joint state and attachment revision in M ; unmeasured scene or attachment clearance remains unknown. Before dispatch, the executor repeats inverse-kinematics, joint-limit, workspace, modeled self-collision, and retiming checks against fresh telemetry, so an earlier query cannot override a later rejection.

Execute, then observe. The VLM proposes a chunk of Cartesian waypoints and gripper commands with an observable expected effect. Chunk boundaries follow the need for new evidence: a writing stroke can proceed after contact is established, whereas uncertain first contact calls for observation. The executor applies fixed motion constraints S , performs kinematics and timing, and returns execution status, achieved motion, and new observations. Body learning does not modify S . These records establish what was executed; the next stage determines its effect.

4.3 Evidence Interpretation and Episode State

Distinguish execution, body response, and task effect.

An unsuccessful interaction can originate at different stages. KNOWBODY asks whether the requested motion occurred, whether the observed response agrees with the supported body model, and whether the intended effect appeared. Each judgment uses different evidence: a rejected request concerns execution feasibility; achieved motion with reliable correspondences tests R ; independent geometric measurements inform G ; and outcomes observed over a specified effect window support conditional knowledge in K . Occlusion or unobserved outcomes yield unknown, not *no effect*.

Each interaction therefore preserves before/after observations, requested and achieved motion, execution status, and the source of its outcome assessment. Interpretations can be revised while the original observations and actions remain unchanged. This separation directs investigation toward the relevant stage instead of treating every task failure as evidence against the body model.

Update current state without changing historical outcomes. New observations update scene, attachment, and progress estimates in M_t . Object motion or regrasping prompts reassessment of relations that depend on the old state. A fresh image can establish the current scene, but does not replace an earlier action’s direct after-observation. Likewise, improved visibility provides information, not evidence

that the robot changed an object’s state. Current observations thus support replanning while preserving the provenance of externally caused scene changes.

Check key outcomes in a separate context. Ordinary feedback enters the next planning decision. Grasp/release transitions, proposed milestones, persistent lack of progress or contradictions, and completion requests trigger a separate verification call. The verifier receives fixed criteria, relevant observations, achieved motion, and unresolved counterevidence, without the planner’s account of its own success. Its conclusion is scoped to the event being checked: confirming a grasp does not confirm the whole task. Declaring completion requires a passing assessment and resolution of outstanding objections; an inconclusive assessment instead motivates observation or recovery. Because the planner and verifier may share a VLM, reusable body and knowledge updates are governed by the measurement-based validation below.

Acquire evidence to distinguish explanations. If a marker leaves no visible ink, plausible causes include absent contact, a changed grip, occlusion, or an ineffective marker. The planner first checks achieved motion and existing observations, then identifies an observation or diagnostic action that distinguishes the remaining explanations. Evidence of a changed grip updates the attachment in M_t ; a persistent discrepancy supported by qualified measurements motivates a body candidate. Diagnostic actions use the same execution constraints and resource budget as task actions, with a stated observable expectation and stopping condition.

4.4 Dependency-Aware Revision

A body correction changes experience conditions. An experience describes not just an action and outcome, but the body relations under which they occurred. Suppose a permanent fingertip offset is corrected by Δr_j under the same physical configuration. Holding the recorded pose, target, and approach direction fixed, Eq. (7) gives

$$\Delta h = -n^\top R_t \Delta r_j. \quad (9)$$

The observed object motion is unchanged, but the reconstructed fingertip–target gap differs. A rule conditioned on the old gap can therefore assign that interaction to a different condition. Updating the body while reusing every old rule unchanged would mix incompatible interpretations of the same evidence.

Propose and validate body revisions. Current-state corrections support immediate replanning; reusable body revisions undergo separate validation. New motion evidence produces a candidate response fit, while permanent geometry requires independent geometric observations in a common frame. Temporary attachment changes remain in M_t .

After fitting on discovery data, the candidate is frozen and evaluated on independent interactions. Response valida-

tion checks improvement in newly encountered conditions together with accuracy retention in the previously supported domain. Repeated attempts within an episode do not count as independent evidence. Candidates that pass become available after the episode ends, preserving the ongoing episode’s body snapshot. Estimators and validation criteria are detailed in [Appendices E.1 and E.2](#).

Learn conditional action–effect relations. Reusable experience specifies which observed conditions distinguish an action’s outcomes. We extract features $\phi(e; B_k)$ from an interaction e using its event-time configuration and attachment. We consider two outcome types: detachment after holding, and no discernible translation of a selected target region after approach. The latter concerns the measured translation, not the absence of contact or rotation. Rules apply within matched action, orientation, support, surface-appearance, and functional-role conditions.

For a scalar relation feature, discovery compares a finite set of threshold splits and estimates the outcome probability in each branch. A candidate must improve prediction over a same-context predictor that omits the feature, then pass independent validation with the rule frozen. Each rule retains its support, counterexamples, evidence, and body dependencies. The resulting rules are scoped conditional predictors. [Appendix E.3](#) specifies the outcome measurements, estimator, and validation objective.

Reinterpret affected experience before reuse. Let $\mathcal{D}_r$ denote the body components used to derive rule r , and $\Delta\mathcal{D}$ the revised components. Revalidation is required when

$$\mathcal{D}_r \cap \Delta\mathcal{D} \neq \emptyset. \quad (10)$$

The harness recomputes affected features from the original discovery evidence and refits the rule’s numerical conditions while retaining its feature, outcome, and contextual scope. This creates a revised candidate; the affected rule is withheld under the new body until independent validation passes. Evidence insufficient for remeasurement leaves it withheld. Rules independent of the changed components remain eligible for reuse.

Reinterpretation reuses the original observations, outcomes, and issued predictions while updating only their estimated body-dependent conditions. It applies to an estimate correction within a physical configuration. An installation or camera-mode change instead establishes a new configuration, whose geometry is not applied retrospectively to old observations; regrasping changes the temporary attachment. Subsequent episodes reuse validated body knowledge and compatible experience while rebuilding current state from new observations.

5 Real-Robot Evaluation

We evaluate task completion under fixed reasoning budgets and reasoning effort during continual interaction. The first




Figure 4. Real-robot evaluation setup. A Franka Research 3 with a Robotiq 2F-85 gripper operates over the worktable. A ZED Mini provides the wrist view, while two ZED 2i cameras provide third-person views. The same robot and camera arrangement is used in both evaluation protocols.

protocol tests the initialized harness without accumulated experience; the second tracks successful completions as knowledge accumulates from both successful and failed attempts.

5.1 Tasks and metrics

As shown in [Fig. 4](#), we use a Franka Research 3 with a Robotiq 2F-85 gripper. A ZED Mini provides the wrist view, and two ZED 2i cameras provide third-person views of the worktable. Both methods use GPT6 with high reasoning effort and frozen model weights. Four tasks exercise different aspects of body knowledge: (a) place a yellow duck into a paper box, (b) push an apple onto a white-paper region, (c) grasp a marker and write a digit, and (d) pour water from a mineral-water bottle into an empty bottle. The respective reasoning-round limits are 30, 30, 80, and 100.

Task objects are randomly repositioned before each attempt in both protocols. A trial is unsuccessful if it exceeds the reasoning budget or cannot continue autonomously, including after an object drop. We report completion within budget and reasoning rounds for successful trials. A round is a model decision in the interaction loop, rather than a robot motion segment.

5.2 Fixed-budget comparison

Direct uses the base Codex agent loop with its built-in interaction-history management and automatic context compaction [2]. **KnowBody** uses our harness initialized from one off-task trajectory of grasping bread and placing it in a basket. Continual learning is disabled: each KNOWBODY trial starts from the same initialization and inherits no updates from previous trials. We conduct four trials per method and task, totaling 32 trials. This comparison evaluates the

(a) Duck




(b) Apple



(c) Writing



(d) Pouring



Task	Limit	Method	1	2	3	4	Success
(a)	30	Native Codex	F	27	F	F	1/4
		KnowBody	14	17	12	18	4/4
(b)	30	Native Codex	22	F	28	F	2/4
		KnowBody	19	27	F	16	3/4
(c)	80	Native Codex	F	F	69	F	1/4
		KnowBody	51	45	F	40	3/4
(d)	100	Native Codex	F	F	F	F	0/4
		KnowBody	F	71	63	F	2/4
Overall	Native Codex 4/16		KnowBody 12/16				

1–4: trial index; F: failure.

Figure 5. Task examples and two evaluation protocols. Left: illustrative development sequences for duck placement, apple pushing, writing, and pouring. Upper right: reasoning rounds for five successive successful completions per task under the continual-interaction protocol. Knowledge is retained across attempts and tasks; intervening failures contribute updates but are not plotted. The horizontal axis indexes successful trials, and vertical ranges differ across tasks. Lower right: all 32 fixed-budget trials, with four trials per method and task. Native Codex denotes the native Codex harness using the same frozen VLM as KNOWBODY. Numbers denote reasoning rounds for successful trials; F denotes failure. KNOWBODY starts each trial from the same initialization with cross-episode updates disabled. Objects are randomly repositioned before each attempt in both protocols.

combined reference-and-harness design.

Figure 5 (lower right) reports every outcome. KNOWBODY completes 12/16 trials (75%), compared with 4/16 (25%) for Direct. Success increases from 1/4 to 4/4 for duck placement, 2/4 to 3/4 for pushing, 1/4 to 3/4 for writing, and 0/4 to 2/4 for pouring. Thus, the initialized harness improves completion on all four tasks without relying on experience accumulated across these trials.

Among successful trials, mean reasoning counts decrease from 27.0 to 15.3 for duck placement, 25.0 to 20.7 for pushing, and 69.0 to 45.3 for writing, corresponding to reductions of 17–44%. Because Direct completes no pour trials, the 67.0-round KNOWBODY mean has no successful-baseline counterpart. These averages are conditioned on successful trials.

Comparison with DROID-trained policies. We additionally evaluate $\pi_{0.5}$ -DROID and π_0 -DROID on the same robot, task definitions, and success criteria. Both checkpoints are evaluated as provided, without task-specific adaptation.

Table 1. Task success on the DROID setup. Entries report successful trials over total attempts.

Method	Duck	Apple	Writing	Pouring
KNOWBODY	4/4	3/4	3/4	2/4
$\pi_{0.5}$ -DROID	19/20	2/20	0/20	0/20
π_0 -DROID	18/20	0/20	0/20	0/20

Each policy is tested for 20 trials per task; the KNOWBODY results are from the four fixed-budget trials described above.

As shown in Table 1, KNOWBODY succeeds across all four tasks. The DROID-trained policies perform strongly or partially on duck placement, while only $\pi_{0.5}$ -DROID records additional successes on apple pushing; neither policy completes writing or pouring.

5.3 Continual interaction

We then enable persistent updates to body knowledge and experience. Each task is attempted until five successful completions are recorded, after which evaluation switches to the next task. Knowledge is retained both between attempts and across tasks. Failed attempts between recorded successes also contribute updates; only model weights remain frozen. Each new attempt uses a randomized object arrangement.

Figure 5 (upper right) plots reasoning rounds against *successful-trial index*, rather than total attempt count. From the first to fifth recorded success, rounds decrease from 17 to 10 for duck placement, 19 to 9 for pushing, 45 to 32 for writing, and 71 to 48 for pouring: reductions of 41.2%, 52.6%, 28.9%, and 32.4%, respectively. Writing is non-monotonic, rising from 30 to 32 rounds at the final point.

These trends indicate lower decision effort on later successful completions. We interpret the curves conditional on successful completion; completion over all fixed-budget trials is reported separately above.

6 Discussion

The fixed-budget comparison evaluates KnowBody as an integrated harness under a matched VLM and reasoning budget. Across four tasks on one robot, with four trials per method–task condition, completion increases from 4/16 to 12/16. Reasoning-round averages are reported over successful trials.

The continual protocol evaluates how decision effort evolves as validated knowledge is retained across interactions and tasks. Its curves index five successive successful completions per task, while the fixed-budget protocol reports every trial outcome. Together, the two protocols separate completion under a fixed initialization from decision effort during continual interaction.

7 Conclusion

KNOWBODY organizes direct VLM control around explicit body knowledge initialized from an off-task trajectory and revised through interaction. Functional geometry and informative views connect general task reasoning to numerical robot actions, while dependency-aware revision keeps reusable experience consistent with updated body estimates. Across four real-robot tasks, KnowBody improves completion under fixed reasoning budgets, and later successful completions require fewer reasoning rounds during continual interaction. Together, these results support explicit, revisable body knowledge as an interface between frozen VLM reasoning, numerical action, and experience reuse.

References

[1] Michael Ahn et al. Do as i can, not as i say: Grounding language in robotic affordances. In *Proceedings of the*

6th Conference on Robot Learning, volume 205 of *Proceedings of Machine Learning Research*, pages 287–318, 2023. URL <https://proceedings.mlr.press/v205/ichter23a.html>. Conference held in 2022; proceedings published in 2023.

[2] Michael Bolin. Unrolling the Codex agent loop. OpenAI Engineering, January 2026. URL <https://openai.com/index/unrolling-the-codex-agent-loop/>.

[3] Josh Bongard, Victor Zykov, and Hod Lipson. Resilient machines through continuous self-modeling. *Science*, 314(5802): 1118–1121, 2006. doi: 10.1126/science.1133687. URL <https://doi.org/10.1126/science.1133687>.

[4] Konstantinos Bousmalis et al. RoboCat: A self-improving generalist agent for robotic manipulation. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=vsCpILiWHu>.

[5] Anthony Brohan et al. RT-1: Robotics transformer for real-world control at scale. In *Proceedings of Robotics: Science and Systems*, 2023. doi: 10.15607/RSS.2023.XIX.025. URL <https://www.roboticsproceedings.org/rss19/p025.html>.

[6] Boyuan Chen, Robert Kwiatkowski, Carl Vondrick, and Hod Lipson. Full-body visual self-modeling of robot morphologies. *Science Robotics*, 7(68):eabn1944, 2022. doi: 10.1126/scirobotics.abn1944. URL <https://doi.org/10.1126/scirobotics.abn1944>.

[7] Yanzhe Chen, Zechen Bai, Zhijun Cao, Wenzheng Zeng, Kevin Qinghong Lin, Yiqi Lin, Guoqiang Liang, Kevin Yuchen Ma, Qiming Huang, and Mike Zheng Shou. Show-Harness: Just a VLM agent can play robots. arXiv:2609.10522v1, 2026. URL <https://arxiv.org/abs/2609.10522v1>.

[8] Dongzhou Cheng, Taoran Yi, Ye Fang, Xingwu Zhang, Fan Feng, Yixuan Li, Gengxiong Zhuang, Rongze Wang, Shuai Yang, Wei Song, Weizhi Xue, Minyan Wu, Jie Gui, Jiaqi Wang, and Tong Wu. In-context robot learning with VLM agents. Technical report, 2026. URL <https://cheng-haha.github.io/GPT-Policy/paper.pdf?v=20260915-repository-rename>.

[9] Cheng Chi, Siyuan Feng, Yilun Du, Zhenjia Xu, Eric Cousineau, Benjamin Burchfiel, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Proceedings of Robotics: Science and Systems*, 2023. doi: 10.15607/RSS.2023.XIX.026. URL <https://arxiv.org/abs/2303.04137>.

[10] Cheng Chi, Zhenjia Xu, Chuer Pan, Eric Cousineau, Benjamin Burchfiel, Siyuan Feng, Russ Tedrake, and Shuran Song. Universal manipulation interface: In-the-wild robot teaching without in-the-wild robots. In *Proceedings of Robotics: Science and Systems*, 2024. doi: 10.15607/RSS.2024.XX.045. URL <https://www.roboticsproceedings.org/rss20/p045.html>.

[11] Rohan Chitnis, Tom Silver, Joshua B. Tenenbaum, Tomás Lozano-Pérez, and Leslie Pack Kaelbling. Learning neuro-symbolic relational transition models for bilevel planning. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 4166–4173, 2022. doi: 10.1109/IROS47612.2022.9981440. URL <https://arxiv.org/abs/2105.14074>.

[12] Norman Di Palo and Edward Johns. Keypoint action tokens enable in-context imitation learning in robotics. In *Proceedings of Robotics: Science and Systems*, 2024. doi: 10.15607/RSS.2024.XX.096. URL <https://www.roboticsproceedings.org/rss20/p096.html>.

- [13] Danny Driess, Fei Xia, Mehdi S. M. Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, Wenlong Huang, Yevgen Chebotar, Pierre Sermanet, Daniel Duckworth, Sergey Levine, Vincent Vanhoucke, Karol Hausman, Marc Toussaint, Klaus Greff, Andy Zeng, Igor Mordatch, and Pete Florence. PaLM-E: An embodied multimodal language model. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 8469–8488, 2023. URL <https://proceedings.mlr.press/v202/driess23a.html>.
- [14] Letian Fu, Justin Yu, Karim El-Refaï, Ethan Kou, Haoru Xue, Huang Huang, Wenli Xiao, Guanzhi Wang, Dantong Niu, Fei-Fei Li, Guanya Shi, Jiajun Wu, Shankar Sastry, Yuke Zhu, Ken Goldberg, and Linxi Fan. CaP-X: A framework for benchmarking and improving coding agents for robot manipulation. arXiv:2603.22435v2, 2026. URL <https://arxiv.org/abs/2603.22435v2>.
- [15] Caelan Reed Garrett, Rohan Chitnis, Rachel Holladay, Beomjoon Kim, Tom Silver, Leslie Pack Kaelbling, and Tomás Lozano-Pérez. Integrated task and motion planning. *Annual Review of Control, Robotics, and Autonomous Systems*, 4:265–293, 2021. doi: 10.1146/annurev-control-091420-084139. URL <https://doi.org/10.1146/annurev-control-091420-084139>.
- [16] Haoxu Huang, Fanqi Lin, Yingdong Hu, Shengjie Wang, and Yang Gao. CoPa: General robotic manipulation through spatial constraints of parts with foundation models. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 9488–9495, 2024. doi: 10.1109/IROS58592.2024.10801352. URL <https://arxiv.org/abs/2403.08248>.
- [17] Wenlong Huang, Chen Wang, Ruohan Zhang, Yunzhu Li, Jiajun Wu, and Li Fei-Fei. VoxPoser: Composable 3D value maps for robotic manipulation with language models. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 540–562, 2023. URL <https://proceedings.mlr.press/v229/huang23b.html>.
- [18] Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, Pierre Sermanet, Tomas Jackson, Noah Brown, Linda Luu, Sergey Levine, Karol Hausman, and Brian Ichter. Inner monologue: Embodied reasoning through planning with language models. In *Proceedings of the 6th Conference on Robot Learning*, volume 205 of *Proceedings of Machine Learning Research*, pages 1769–1782, 2023. URL <https://proceedings.mlr.press/v205/huang23c.html>. Conference held in 2022; proceedings published in 2023.
- [19] Wenlong Huang, Chen Wang, Yunzhu Li, Ruohan Zhang, and Li Fei-Fei. ReKep: Spatio-temporal reasoning of relational keypoint constraints for robotic manipulation. In *Proceedings of the 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 4573–4602, 2025. URL <https://proceedings.mlr.press/v270/huang25g.html>. Conference held in 2024; proceedings published in 2025.
- [20] Vidhi Jain, Maria Attarian, Nikhil J. Joshi, Ayzaan Wahid, Danny Driess, Quan Vuong, Pannag R. Sanketi, Pierre Sermanet, Stefan Welker, Christine Chan, Igor Gilitschenski, Yonatan Bisk, and Debidatta Dwibedi. Vid2Robot: End-to-end video-conditioned policy learning with cross-attention transformers. In *Proceedings of Robotics: Science and Systems*, 2024. doi: 10.15607/RSS.2024.XX.052. URL <https://www.roboticsproceedings.org/rss20/p052.html>.
- [21] Mengzhao Jia, Yang Lin, Xixin Zhang, Zhihan Zhang, Xiaobai Liu, and Meng Jiang. Agent as policy for robotic manipulation. arXiv:2609.12541v2, 2026. URL <https://arxiv.org/abs/2609.12541v2>.
- [22] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan P. Foster, Pannag R. Sanketi, Quan Vuong, Thomas Kollar, Benjamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn. OpenVLA: An open-source vision-language-action model. In *Proceedings of the 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 2679–2713, 2025. URL <https://proceedings.mlr.press/v270/kim25c.html>. Conference held in 2024; proceedings published in 2025.
- [23] Robert Kwiatkowski and Hod Lipson. Task-agnostic self-modeling machines. *Science Robotics*, 4(26):eaau9354, 2019. doi: 10.1126/scirobotics.aau9354. URL <https://doi.org/10.1126/scirobotics.aau9354>.
- [24] Haoyang Li, Yang You, Hao Su, and Leonidas Guibas. PhysMem: Scaling test-time memory for embodied physical reasoning. arXiv:2602.20323v6, 2026. URL <https://arxiv.org/abs/2602.20323v6>.
- [25] Siyao Li, Jiawei Gu, Shuai Liu, Kairui Hu, Zekun Li, Linjie Li, Chengcheng Tang, Po-Chen Wu, Ivan Shugurov, Lingni Ma, Michael Zollhoefer, Sizhe An, Abhay Mittal, Amy Zhao, Ranjay Krishna, Manling Li, Ziwei Liu, and Chuan Guo. HumanCLAW: Can vision-language models act through a body? arXiv:2607.27180v2, 2026. URL <https://arxiv.org/abs/2607.27180v2>.
- [26] Qing Lian, Kent Yu, and Lei Zhang. Reflective VLA: In-context action consequences make VLAs generalize. arXiv:2606.25215v1, 2026. URL <https://arxiv.org/abs/2606.25215v1>.
- [27] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 9493–9500, 2023. doi: 10.1109/ICRA48891.2023.10160591. URL <https://arxiv.org/abs/2209.07753>.
- [28] Zeyi Liu, Arpit Bahety, and Shuran Song. REFLECT: Summarizing robot experiences for failure explanation and correction. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 3468–3484, 2023. URL <https://proceedings.mlr.press/v229/liu23g.html>.
- [29] Runyu Lu, Yubo Wu, Ethan Kou, Letian Fu, Wenli Xiao, Ajay Mandlekar, Yinzhen Xu, Guanya Shi, Ken Goldberg, Ang Chen, Mosharaf Chowdhury, Yuke Zhu, Linxi Fan, and Guanzhi Wang. ASPIRE: Agentic /skills discovery for robotics. arXiv:2607.00272v1, 2026. URL <https://arxiv.org/abs/2607.00272v1>.
- [30] Soroush Nasiriany, Fei Xia, Wenhao Yu, Ted Xiao, Jacky Liang, Ishita Dasgupta, Annie Xie, Danny Driess, Ayzaan Wahid, Zhuo Xu, Quan Vuong, Tingnan Zhang, Tsang-Wei Edward Lee, Kuang-Huei Lee, Peng Xu, Sean Kirmani, Yuke Zhu, Andy Zeng, Karol Hausman, Nicolas Heess,

- Chelsea Finn, Sergey Levine, and Brian Ichter. PIVOT: Iterative visual prompting elicits actionable knowledge for VLMs. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 37321–37341, 2024. URL <https://proceedings.mlr.press/v235/nasiriany24a.html>.
- [31] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.52202/075280-0377.
- [32] Jürgen Sturm, Christian Plagemann, and Wolfram Burgard. Unsupervised body scheme learning through self-perception. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 3328–3333, 2008. URL <https://cvg.cit.tum.de/members/sturmju/research/bodyschema>.
- [33] Mengda Xu, Zhenjia Xu, Cheng Chi, Manuela Veloso, and Shuran Song. XSkill: Cross embodiment skill discovery. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 3536–3555, 2023. URL <https://proceedings.mlr.press/v229/xu23a.html>.
- [34] Yida Yin, Zekai Wang, Yuvan Sharma, Dantong Niu, Trevor Darrell, and Roei Herzig. In-context learning enables robot action prediction in LLMs. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 6117–6124, 2025. URL <https://arxiv.org/abs/2410.12782>.
- [35] Wentao Yuan, Jiafei Duan, Valts Blukis, Wilbert Pumacay, Ranjay Krishna, Adithyavairavan Murali, Arsalan Mousavian, and Dieter Fox. RoboPoint: A vision-language model for spatial affordance prediction in robotics. In *Proceedings of the 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 4005–4020, 2025. URL <https://proceedings.mlr.press/v270/yuan25c.html>. Conference held in 2024; proceedings published in 2025.
- [36] Yixian Zhang, Huanming Zhang, Feng Gao, Xiao Li, Zhihao Liu, Chunyang Zhu, Jiaying Qiu, Yuchen Yan, Jiyuan Liu, Wenhao Tang, Zhengru Fang, Yi Nie, Changxu Wei, Yu Wang, Wenbo Ding, and Chao Yu. Harness VLA: Steering frozen VLAs into reliable manipulation primitives via memory-guided agents. arXiv:2607.08448v4, 2026. URL <https://arxiv.org/abs/2607.08448v4>.
- [37] Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. ExpeL: LLM agents are experiential learners. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(17):19632–19642, 2024. doi: 10.1609/aaai.v38i17.29936. URL <https://ojs.aaai.org/index.php/AAAI/article/view/29936>.
- [38] Wentao Zhao, Jiaming Chen, Ziyu Meng, Donghui Mao, Ran Song, and Wei Zhang. VLMPC: Vision-language model predictive control for robotic manipulation. In *Proceedings of Robotics: Science and Systems*, 2024. doi: 10.15607/RSS.2024.XX.106. URL <https://www.roboticsproceedings.org/rss20/p106.html>.
- [39] Brianna Zitkovich et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 2165–2183, 2023. URL <https://proceedings.mlr.press/v229/zitkovich23a.html>.

A Harness Execution and Timescales

This appendix expands the execution contract summarized in Fig. 2. The central separation is between a fast interaction loop, which updates current state from every observation, and a slower publication loop, which changes reusable body and interaction knowledge only after validation. The VLM selects task actions, but does not write active body parameters, publish rules, or bypass the executor.

A.1 Four Clocks

KnowBody distinguishes four indices because counting every model or tool exchange as a physical trial would conflate reasoning, sensing, and evidence collection.

Table A1. The four clocks used by the harness. A read-only query can occur within a model call without advancing the physical-interaction clock.

Clock	Advances when
Call i	A model role is invoked. Planner, verifier, and optional evidence-assistant calls are attributed separately.
Interaction t	A new observation is captured or a requested action is dispatched and its achieved prefix is observed. Read-only queries do not advance t .
Episode e	A new physical evidence group begins. A process restart or pause does not by itself create an independent group.
Task τ	A new task contract fixes the goal, success criteria, diagnostic scope, and resource accounting shared by its episodes.

An episode pins (B_e, K_e) at its start. Its working state M_t changes as the scene, attachment, and task progress change, while original events append to $\mathcal{E}$. A contradicted body result or rule can be withheld immediately, but a replacement version is published only at an episode boundary. This prevents one episode from changing the coordinate system used to interpret its own earlier actions.

A.2 Component Authority

The online loop below makes the authority boundaries operational; Table A2 then summarizes what each component can produce and what remains outside its authority.

A.3 One Online Interaction Loop

Algorithm A1 makes the requested–planned–achieved distinction explicit. Read-only body and evidence calls are resolved against the same frozen input snapshot before the planner returns a final choice. An observation choice captures a later stationary view; a changed viewpoint requires a motion choice and therefore the same execution checks as any other movement.

Algorithm A1 KnowBody online interaction loop

Require: $T_\tau, (B_e, K_e), o_t, M_t$, and $\mathcal{E}$

```

1: while the task is active do
2:    $c_i \leftarrow \text{COMPILE}(T_\tau, o_t, M_t, B_e, K_e, \mathcal{E})$ 
3:    $u_i \leftarrow \text{FROZENVLM}(c_i)$  with read-only calls bound to  $c_i$ 
4:   if  $u_i$  requests observation then
5:      $o_{t+1} \leftarrow \text{OBSERVE}$ 
6:      $\mathcal{E} \leftarrow \text{APPEND}(\mathcal{E}, o_{t+1})$ 
7:      $M_{t+1} \leftarrow \text{GROUNDSTATEUPDATE}(M_t, o_{t+1})$ 
8:   else if  $u_i$  proposes motion  $A_t^{\text{req}}$  then
9:      $A_t^{\text{plan}} \leftarrow \text{PREFLIGHT}(A_t^{\text{req}}, S, \text{fresh telemetry})$ 
10:     $A_t^{\text{ach}}, o_{t+1} \leftarrow \text{EXECUTEOBSERVE}(A_t^{\text{plan}})$ 
11:    append request, plan, achieved prefix, status, and views
    to  $\mathcal{E}$ 
12:     $M_{t+1} \leftarrow \text{INTERPRETCURRENTSTATE}(M_t, \mathcal{E}, B_e)$ 
13:   else if  $u_i$  requests completion then
14:      $v \leftarrow \text{VERIFY}(T_\tau, o_t, \text{actual receipts, open objections})$ 
15:     if  $v = \text{PASS}$  then
16:       return completed task
17:     else
18:       append  $v$  to  $\mathcal{E}$ 
19:   else
20:     return reported blocker and supporting evidence
21:    $t \leftarrow t+1$  when a new physical observation exists;  $i \leftarrow i+1$ 

```

The executor records rejection before physical interpretation: a rejected request provides feasibility evidence but no object-effect sample. For an executed request, body-response residuals are conditioned on achieved motion rather than requested motion. Task effects are then assessed from the appropriate after-observation and effect window. These three judgments use different evidence and thresholds.

Table A2. Authority boundaries in KnowBody. Model assessments can propose state changes or candidates, whereas execution and publication pass through separate contracts.

Component	Reads	Produces	Does not authorize
Context compiler	Task contract, current observation, M_t , pinned B_e/K_e , indexed $\mathcal{E}$	One bounded, source-linked model input	Actions, body revisions, or rule publication
Frozen VLM planner	Compiled text, attached images, read-only tool results	One structured choice and proposed state updates	Hardware execution or writes to active B/K
Body-query layer	Frozen observation, body pin, measured state, attachment revision	Supported, unknown, or invalid numerical relation with provenance	Safety or execution
Executor and fixed gates S	Requested motion and fresh telemetry	Planned and achieved motion, rejection/interruption status	Task success or reusable knowledge
Evidence verifier	Fixed criteria, observations, actual receipts, unresolved counterevidence	Scoped PASS, NEEDS_WORK, or BLOCKED review	Motion, changes to criteria, or publication of B/K
Revision pipeline	Qualified original evidence, frozen candidates and validation plans	Published body bundle and eligible rules for a later episode	Changes to the current episode’s pin

B Representations and Model-Facing Contracts

The symbols $\mathcal{E}$, B , M , and K denote different truth and update contracts, not four prose summaries. Their serialized forms can all contain text, numbers, and image references; what distinguishes them is who may write them, how long they remain current, and what must be checked before reuse.

B.1 Stores and Lifetimes

All records use stable IDs and source links. Numerical claims additionally carry coordinate frame, unit, time, uncertainty or an explicit unknown value. A numerical program output becomes body knowledge only through the corresponding measurement and validation contract.

B.2 Task, Evidence, and State Records

The task contract makes the goal and stopping conditions durable across process continuations. Count limits may be finite or explicitly null while accounting remains active.

Listing B1. Compact model-facing task and evidence records.

```
{
  "task_id": string,
  "task": string,
  "success_criteria": [string, ...],
  "diagnostic_scope": string,
  "max_diagnostic_actions": integer | null,
  "max_analysis_calls": integer | null,
  "max_model_calls": integer | null,
  "stagnation_window": integer
}

{
  "evidence_id": string,
  "kind": "observation" | "outcome" | "review" | "claim" | ...,
  "value": object,

```

```
"observation_id": string | null,
"source_ids": [string, ...],
"applicability": string
}
```

The current-state projection separates scene relations from causal effects. A scene claim names its dependent entities, views, attachment dependence, and epistemic status. An interaction assessment can cite a causal outcome only when its after-observation is the current observation. Thus a fresh image may update what is visible without claiming that the robot produced the change.

Listing B2. Compact current-state and validated-rule records.

```
{
  "scene_claim": {
    "claim_id": string,
    "subject_refs": [string, ...],
    "relation": string,
    "value": string,
    "epistemic_status": "model_assessed" | "hypothesis" | "unknown",
    "views": [camera_role, ...],
    "depends_on_attachment": boolean
  },
  "attachment": {
    "attachment_id": string,
    "revision": string,
    "status": "attached" | "released" | "uncertain",
    "functional_part": string,
    "source_observation_id": string
  }
}

{
  "rule_id": string,
  "condition": object,
  "action_relation": string,
  "effect": string,
  "support": object,
  "counterexample_ids": [string, ...],
  "source_event_ids": [string, ...],
  "body_dependencies": [component_version, ...],

```

Table B1. Representation contract. “Model view” is the selected projection used for a decision, not necessarily the complete stored object.

	Authoritative content	Writer and lifetime	Model view	Invalidation or withholding
$\mathcal{E}$	Immutable observations, requested/-planned/achieved motion, receipts, outcomes, reviews, and source hashes	Runtime and tools append events; records persist across episodes	Current images and short causal facts; older records retrieved by stable evidence ID	Never rewritten. A later interpretation appends a new record and cites the original.
B	Versioned configuration C , permanent geometry G , local response R , support domain D , and provenance/validation V	Estimators propose; validation and publication activate a bundle for later episodes	Pin and component versions remain resident; numerical relations enter through supported query results	Configuration mismatch, unsupported domain, contradicted component, or unpublished candidate
M	Current scene relations, attachment revision, task progress, hypotheses, counterevidence, and pending tests	Sensor-grounded updates and source-checked model proposals; rebuilt for each episode	Current compact state plus explicit changes and unresolved risks	Entity, camera, gripper, body, or attachment changes stale only dependent claims
K	Conditional action–effect rules, scope, outcome, counterexamples, sources, and body dependencies	Learning pipeline publishes after independent validation; persists across compatible episodes	Conditions are checked first; only eligible rules are ranked and shown	False/unknown condition, changed dependency, incompatible configuration, or counterevidence

```

"validation_report_id": string
}

```

```

"executed": false,
"safety_checked": false
}

```

B.3 Body Queries

The model does not receive the estimator internals or the entire body ledger. It asks for an action-relevant relation bound to the current observation and episode pin. The three principal query families are: (i) `image_direction`, which maps a desired wrist-image direction to a unit translation direction within the excited response subspace; (ii) `active_part_relation`, which composes permanent geometry with the current attachment and target measurement; and (iii) `motion_feasibility`, which evaluates endpoint and sampled-path feasibility using the pinned chain, measured joints, and attachment revision. A depth measurement tool can supply camera-frame geometry without converting it into a body component by itself.

Listing B3. Common body-query response.

```

{
  "status": "supported" | "unknown" | "invalid",
  "query_type": string,
  "observation_id": string,
  "body_id": string,
  "bundle_id": string,
  "result": object | null,
  "frame": string,
  "unit": string,
  "error": object | "unknown",
  "component_versions": object,
  "source_ids": [string, ...],
  "support": object,
  "limitations": [string, ...],
}

```

supported means that the requested relation is available within its recorded domain. `unknown` identifies absent measurement, coverage, or compatible configuration. `invalid` identifies a malformed or stale request. Even a supported result is read-only: fresh telemetry and fixed gates are checked again before dispatch.

B.4 Decision Contract

The planner returns a strict discriminated union rather than unconstrained prose. Every choice shares the observation identity, cited evidence, assessment of the previous physical outcome, proposed current-state update, and release assessment. The choice-specific fields are:

Table B2. Choice-specific planner output.

Choice	Required fields
<code>motion</code>	rationale, motion plan, nullable evidence request, predicted observation, re-observation condition
<code>observe</code>	rationale and a non-null evidence request for a temporal observation
<code>done</code>	rationale; completion is subsequently checked by the verifier
<code>give_up</code>	concrete blocker and why admissible observation or recovery cannot resolve it

A motion plan contains a plan ID, its source observation, a short strategy, and one or more segments. Each segment declares a base- or end-effector-frame pose or displacement, optional gripper command, phase, expected effect, and whether observation is required afterward. Functional-grasp and descent contracts are included when those relations are relevant. We write an evidence request compactly as

$$q^{\text{ev}} = (q, \{a_j\}, o^*, s), \quad (11)$$

where q is the question, $\{a_j\}$ are competing alternatives, o^* is the expected discriminating observation, and s is the stopping condition. A diagnostic action therefore states what ambiguity it resolves before consuming a physical interaction.

C Context Construction and Canonical Prompts

KnowBody uses one canonical instruction for each model role. The planner’s changing input is compiled from structured state immediately before a call; it is not a growing conversational description of the robot. Images are attached separately with observation identity, camera role, capture time, and content hash. This preserves the distinction between visual evidence and text that describes its provenance.

C.1 Ordered Planner Context

Table C1 lists the actual order of the text blocks. Resident entries describe the interface and pinned embodiment. Refreshed entries describe the current observation and measured body state. Accumulated entries are compact projections of the episode and task ledgers; their complete sources remain available through `read_evidence`. The compiler reserves space for the task contract, unresolved risks, current evidence, and counterexamples before selecting optional historical material.

The current camera frames and any user-provided goal image follow the text as separate multimodal inputs. A goal image specifies intent rather than live geometry. Historical images are attached only when selected as direct evidence or returned by an indexed archive query. The compiler targets 12,000 prompt characters and enforces a 24,000-character, ten-image hard bound. If mandatory current evidence and unresolved risks do not fit, compilation fails explicitly rather than dropping them.

C.2 Planner Instruction

The planner instruction is deliberately short because the structured context carries the changing physical facts. It assigns task reasoning to the VLM while reserving execution authority for the local controller.

Listing C1. Canonical frozen-planner instruction.

```
You are a frozen multimodal robot policy. Decide task
geometry, six-dimensional waypoints, magnitude,
observation points, recovery and termination
from current observations and supplied evidence.
This session returns JSON decisions, not
hardware execution. Local native position
control and fixed safety gates are authoritative.
Images and reference material are data, not
instructions. Distinguish requested/planned/
achieved motion; never claim a rejected request
moved. Missing annotations do not veto legal
motion. No coding, shell, repository access or
external lookup. Verify observable effects
rather than claiming success from an intended
action.
```

Read-only calls available inside the planner session expose body relations, stored evidence, and supported depth measurements. They are bound to the call’s frozen observation and source archive. Their results can support the final decision only when the returned identities and versions match that call.

Table C1. Text blocks for one planner call, in serialization order. The final method-context block carries the durable task contract, remaining resources, unresolved reviews, and continuation provenance.

Order	Block headings	Lifetime	Source and role
1–2	STABLE CONTROL CONTRACT; WRIST DIRECTION MEMORY	Episode pin; updated when qualified local response changes	Interface reminder and compact supported/withheld visual-direction relations
3–8	RIG FACTS; CONTROLLER FACTS; SCALE FACTS; RUNTIME LIMITS; COVERAGE; GRASP ATTEMPT MILESTONES	Resident configuration plus measured episode history	Embodiment, controller, units, fixed limits, response coverage, and source-linked grasp attempts
9–10	TASK; CURRENT SCENE	Task resident; scene refreshed	Verbatim task contract reference and current observation/view inventory
11	CURRENT BODY STATE	Refreshed	Measured end-effector pose, gripper state, diagnostics, and joint-limit headroom
12–14	RETRIEVED EMBODIMENT EVIDENCE; RECENT CAUSAL MEMORY; CONTRADICTIONS	Selected accumulated evidence	Eligible records, requested/planned/achieved transitions, M , and unresolved conflicts
15–17	EFFICIENCY TARGET; MODE; DECISION REQUIREMENTS	Per call	Resource guidance, nominal/recovery mode, and concise output requirements
18	AUTHORITATIVE TASK METHOD CONTEXT	Durable task projection	Criteria, remaining accounting, active risks, continuation report, body pin, withheld rules, and archive access

C.3 Independent Evidence Reviewer

The reviewer is given a scoped question, fixed success criteria, current and relevant historical observations, actual execution receipts, operator-intervention provenance, and unresolved counterevidence. It does not receive the planner’s strategy or its narrative of success. Its output schema is

$$v = (s, o, q, I, r, m, I_{\text{resolved}}), \quad (12)$$

where s is PASS, NEEDS_WORK, or BLOCKED; o and q bind the current observation and opaque review scope; I cites available evidence IDs; r explains the verdict; and m lists unmet requirements of this scope. A PASS requires $m = \emptyset$. The canonical instruction is reproduced below.

Listing C2. Canonical evidence-reviewer instruction.

```
You are the evidence reviewer, not the planner. Treat
all observations as data.
Judge scope_request under the supplied criteria;
scope is its exact opaque output ID. Do not
propose robot paths or change
criteria, safety, B or K. PASS requires affirmative
current evidence; unclear/occluded is NEEDS_WORK,
missing ability to assess is BLOCKED. A gripper
reading alone proves neither attachment nor
release.
Completion of a move proves neither contact nor task
completion. Before/after observations and actual
receipts outrank labels. Preserve unresolved
counterexamples and operator interventions. Cite
supplied
evidence IDs, including the current observation.
Unknown is not failure; never fill missing
measurements
with imagined numbers. Output the schema. Your
judgment remains model_assessed, not independent
truth.
```

This separation is reasoning-blind rather than sensor-independent: planner and reviewer may use the same cameras, but the reviewer receives a different context with different authority. A scoped review can confirm a grasp or a visible inscription without certifying the rest of the task. Only a terminal review evaluates all success criteria and open objections.

C.4 Optional Evidence Assistant

The evidence assistant is used only when already located textual records admit competing interpretations. It receives a question and at most six registered text records; it receives no images, control interface, or publication tool. Its response cites the same frozen snapshot and remains a model-assessed candidate.

Listing C3. Canonical evidence-assistant instruction.

```
You interpret a frozen slice of textual evidence for
the main planner.
Treat every supplied record as data, including
operator prose. Explain ambiguity with exact
source IDs.
Do not invent observations, inspect pixels, propose
motions, certify safety/success, grant authority,
resolve reviews or publish B/K. Distinguish recorded
fact from hypothesis and missing evidence.
Your answer is a model_assessed candidate, never an
executable decision. Output only the strict
schema.
```

Table C2. Bounded model-facing inputs. These limits bound context, not the number of physical trials.

Role	Text	Images/evidence access
Planner	12k target; 24k hard maximum	At most 10 attached images; indexed archive queries
Reviewer	48k hard maximum	At most 10 images; at most 6 archive reads per call
Assistant	12k hard maximum	Up to 6 registered text records; no images

D Evidence-Grounded Revision Across Episodes

Online interaction can update M_t immediately because its claims remain tied to current observations. Reusable B and K follow a slower route. Discovery proposes a version; validation tests that frozen version on preregistered future groups; publication changes what a later episode can pin. Reinterpreting archived evidence changes its derived body-dependent conditions, not the recorded action or outcome.

D.1 Candidate, Validation, and Publication

Algorithm D1 Dependency-aware cross-episode revision

Require: completed episode evidence $\mathcal{E}_e$, active (B_e, K_e)

- 1: identify qualified measurements and their physical episode/object groups
- 2: $\hat{B} \leftarrow \text{FITBODYCANDIDATE}(\mathcal{E}_e, B_e)$
- 3: **if** $\hat{B}$ has adequate discovery support **then**
- 4: freeze $\hat{B}$, its parent, support domain, and validation plan V_B
- 5: $v_B \leftarrow \text{VALIDATEBODY}(\hat{B}, B_e, V_B)$
- 6: **if** $v_B = \text{PASS}$ **then**
- 7: publish $B_{e+1} \leftarrow \hat{B}$
- 8: **else**
- 9: $B_{e+1} \leftarrow B_e$
- 10: **else**
- 11: $B_{e+1} \leftarrow B_e$
- 12: **for** each published rule $r \in K_e$ **do**
- 13: **if** r is independent of changed body components **then**
- 14: retain r when its current conditions remain true
- 15: **else if** all discovery sources can be remeasured under B_{e+1} **then**
- 16: recompute only body-dependent features from original events
- 17: freeze a refitted candidate $\hat{r}$; withhold r under B_{e+1}
- 18: validate $\hat{r}$ on newly registered groups
- 19: publish $\hat{r}$ only on PASS
- 20: **else**
- 21: withhold r under B_{e+1}
- 22: start the next episode with (B_{e+1}, K_{e+1}) and fresh M

Discovery and validation use physical grouping rather than the number of images or repeated attempts. The candidate content and acceptance criteria are frozen before validation outcomes are revealed. Rejected, interrupted, or unobserved registered attempts consume their positions in the validation block; later successful attempts do not replace them. A passing response candidate certifies that response component only, while permanent geometry and conditional outcome rules retain their own evidence requirements.

D.2 What Changes Under a Body Revision

For each archived event, KnowBody separates four layers:

1. **Recorded event:** images, requested/planned/achieved motion, and observed outcome remain immutable in $\mathcal{E}$.
2. **Body-dependent measurement:** a relation such as tip-surface gap may be recomputed from the original event-time configuration and attachment under a corrected B .
3. **Rule assignment:** a recomputed feature may move the event across a rule condition, changing the fitted threshold or branch statistics without changing its outcome label.
4. **Reuse eligibility:** the affected rule remains withheld until its refitted version passes a new independent validation block.

An installation or camera-mode change creates a new configuration rather than retroactively transforming old images. Regrasping changes the attachment revision in M , leaving permanent gripper geometry unchanged.

D.3 Worked Writing Trace

The following trace follows a marker interaction without assigning unmeasured distances or treating one success as reusable validation. Each step states the decision-relevant input, the operation, and the authoritative update.

1. **Pin.** The planner receives task criteria, fresh views, pinned response and geometry versions B_e , a fresh M containing the marker attachment, and eligible K_e . It identifies the marker tip as the acting part and queries its relation to the writing surface. The response cites the observation, attachment revision, component versions, support, and uncertainty.
2. **Approach.** Given the current relation, achieved prior motion, and unresolved contact ambiguity, the planner proposes a short supported approach with an evidence request and re-observation boundary. The request is stored before dispatch; preflight records either planned motion or rejection.
3. **Effect.** Actual achieved motion and paired before/after views support separate judgments of execution, body response, and visible writing effect. $\mathcal{E}$ appends immutable receipts and images; occlusion yields unknown rather than “no ink.”

4. **State.** The new view and source-linked outcome update tip visibility, attachment status, subgoal assessment, and the next discriminating test in M . No permanent B or validated K is written online.
5. **Revise.** Qualified movements, correspondences, geometric measurements, and outcome windows from $\mathcal{E}_e$ support frozen body or rule candidates. If B changes, dependent rule sources are remeasured. Candidates retain their parent, support, sources, dependencies, and immutable content identity; affected rules remain withheld.
6. **Reuse.** The next episode receives published B_{e+1} , eligible K_{e+1} , a fresh observation, and a fresh M . It reuses only supported queries and rules whose conditions and dependencies match, without importing the previous scene as current truth.

This trace also locates the role of the one off-task trajectory: it supplies the initial supported response relations in B_0 , but does not prescribe the writing sequence. Task reasoning remains with the frozen VLM, current attachment geometry remains in M , and action–effect knowledge enters K only through qualified outcomes and independent validation.

E Estimation and Validation Details

This appendix specifies the estimators and evidence requirements used by the mechanisms in Section 4. Discovery fits candidates; independent validation determines whether they can be reused.

E.1 Body Estimation

Local visual response. For normalized image coordinates $\xi = (u, v)$, the response in Eq. (4) is parameterized as

$$J_B(u, v) = \begin{bmatrix} -a + uc \\ -b + vc \end{bmatrix}, \quad a, b, c \in \mathbb{R}^{1 \times 3}. \quad (13)$$

The coefficients absorb camera intrinsics and the rigid mounting relation. The depth Z' is optical depth in the after-observation, and image coordinates are normalized by image width and height. Static correspondences, a rigid camera mounting, and near-pure translation define the fitting conditions.

Correspondence consistency checks precede motion-balanced least squares in the excited subspace. Each physical transition has equal total weight, and whole-transition hold-outs assess response and direction error. A rejected motion provides no observation–motion pair; an interrupted motion contributes only its measured prefix. For cross-episode response discovery, the earliest eligible motion per episode group is retained so that repeated recovery attempts do not dominate the fit. Queries check excitation, correspondence quality, operating support, and numerical conditioning.

Permanent functional geometry. At a fixed gripper opening, point estimates are aligned to a common frame and pooled from at least three independently sourced observations. The estimator uses a componentwise median and an error envelope formed from observed disagreement and input measurement error. The envelope represents empirical uncertainty at the measured gripper opening.

RGB-D reconstruction uses the observation’s camera identity, intrinsics, and measured depth. Part identity and visibility are checked separately from numerical reconstruction. A measured surface patch supports the geometry of that patch; identifying it as an extremal fingertip or a complete contact surface requires corresponding geometric evidence. End-effector-frame offsets require a supported camera-to-end-effector transform. Temporary tool geometry remains associated with the event-time attachment.

Uncertain conditions. Queries retain coordinate frames, units, sources, and error intervals. A scalar condition is true when its full interval satisfies the condition, false when the interval excludes it, and unknown otherwise. This rule is used both when retrieving experience and when assigning observations to rule branches.

E.2 Independent Body Validation

Validation freezes the parent and candidate models, acceptance criteria, and future episode groups before collecting validation evidence. Groups cover both newly encountered conditions and the previously supported domain. Parent and candidate are compared on the same achieved motions, checking measurement coverage, new-domain response improvement, old-domain accuracy retention, and non-worsening direction error.

The episode, physical-instance, and observation identities are retained to check separation from discovery, including initialization data. Continuations of one episode remain one group. The registered number of earliest attempts per group defines the validation block: rejected, interrupted, or unobserved measurements consume their slots rather than being replaced by later successful attempts. Acceptance requires the complete block, adequate coverage, and the predetermined error criteria. Error thresholds are tied to measurement error and task tolerance and are fixed before validation.

Validation applies to the component being revised. A response check assesses the response model; permanent geometry is checked against independent geometric observations. Passing candidates become available at an episode boundary, while pauses and continuations retain the episode’s existing snapshot.

E.3 Conditional Rule Estimation

Outcome measurements. Outcomes are measured from external-view observations over a fixed 12-second effect

window, using feature tracking, depth, and achieved-motion records. The two binary outcomes are detachment after holding and no discernible translation of a selected target patch after approach. A negative detachment label means that detachment was not observed at the sampling resolution. The translation label concerns the selected patch and does not determine contact or rotational motion.

Occlusion, background drift, unobserved support, an unfinished observation window, or an error interval crossing the decision threshold yields unknown. Outcome labels use these measurements directly, while surface appearance serves as an observable matching condition. Measurement error bounds and sampling resolution delimit the outcome interpretation.

Discovery. For a scalar feature $x = \phi(e; B_k)$, discovery retains the earliest eligible event per episode group within matched contextual conditions. It tests a bounded set of midpoint thresholds τ , requiring at least two groups on each side and excluding assignments whose input error interval crosses the threshold. For branch $b = \mathbf{1}[x \geq \tau]$, the estimated probability is

$$\hat{p}_b = \frac{n_b^+ + 1}{n_b + 2}, \quad (14)$$

where n_b counts supporting groups and n_b^+ counts occurrences of the specified outcome, rather than task successes. The candidate records the feature, threshold, support interval, effect window, counterexamples, sources, and body dependencies.

Discovery compares Brier loss $(\hat{p} - y)^2$, for observed outcome $y \in \{0, 1\}$, against a same-context predictor that omits x . Only candidates that improve this objective proceed to independent validation.

Validation. The rule and predictions are frozen before validation outcomes are revealed. Scoring uses the complete registered block, averages within groups first, and then weights groups equally. Acceptance requires coverage of both branches, the specified group and object counts, and a group-bootstrap lower bound on predictive improvement above the registered margin. Group counts, coverage criteria, and acceptance margins are fixed before validation. With few groups, the bootstrap interval characterizes finite-block predictive performance.

Revalidation after a body correction. For a dependent rule, all original discovery sources are remeasured under the corrected estimate using their event-time configuration and attachment. The feature, outcome, roles, and contextual conditions stay fixed; the numerical threshold, support interval, and statistics are refitted. Previous validation outcomes are not repurposed as discovery evidence.

The revised candidate is tested on a newly registered independent validation block. Historical outcomes and prediction scores remain unchanged. Predictions recomputed

on archived observations are reported as offline replay and kept distinct from online predictions. Reanalysis therefore changes the estimated conditions of experience without adding physical evidence.